

# Flights Graph Analysis

February 23, 2024

## 1 Flight Connections and Airports Graph Analysis

```
[ ]: '''  
In this analysis, I am going to use graph analysis applied to a Big Data_  
    ↳framework.  
I will use GraphFrames Spark library with the API for DataFrames (GraphFrames_  
    ↳Spark library)  
to analyze a dataset containing information about flight connections and_  
    ↳airports in the whole world.  
  
I am going to use datasets containing informations about airports, airlines and_  
    ↳flights worldwide.  
Consider these three `csv` files:  
- airports.csv: contains one line for each airport in the world. Among the_  
    ↳others, it provides the columns:  
id, name, city, country, iata, latitude and longitude.  
- airlines.csv: provides some information for each airline. Among the others,_  
    ↳it provides the columns:  
airline_id, name, country, icao.  
- routes.csv enumerates the flights provided by each airline between two_  
    ↳airports. it provides the columns:  
airline_id, airport_source_id, airport_destination_id.  
  
For this Analysis, I am going to use GraphFrames Spark library with the API for_  
    ↳DataFrames  
(GraphFrames Spark library)  
'''
```

```
[ ]: # Input Datasets from the big data cluster
```

```
[2]: airports = '/data/students/bigdata_internet/lab5/airports.csv'  
airlines = '/data/students/bigdata_internet/lab5/airlines.csv'  
flights = '/data/students/bigdata_internet/lab5/routes.csv'
```

```
[ ]: '''In this analysis, PySpark was utilized for its robust distributed computing_  
    ↳capabilities,
```

ideal for handling large datasets efficiently.  
 If you're using the PySpark shell, no additional setup is necessary.  
 However, for those working in a Python environment, setting up PySpark involves  
 ↳ the following steps:

1. Install PySpark: Begin by installing PySpark using pip:  
`pip install pyspark`
2. Configure PySpark.sql: In your Python script or interactive session, include  
 ↳ the following configuration  
 to initialize PySpark.sql:

```
```python
from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()
```
```

Ensure to execute this configuration before performing any PySpark operations.  
 For comprehensive installation and configuration instructions, refer to the  
 ↳ official PySpark documentation:  
[PySpark Installation Guide](#)

```
[ ]: # Reading datasets
```

```
[3]: airportsDF = spark.read.load(airports, format="csv", header=True,
    ↳ inferSchema=True, sep=',')
    airlinesDF = spark.read.load(airlines, format="csv", header=True,
    ↳ inferSchema=True, sep=',')
    flightsDF = spark.read.load(flights, format="csv", header=True,
    ↳ inferSchema=True, sep=',')
```

```
[ ]: # Find top airports and airlines
```

```
[ ]: # Find countries with more than 200 airports
```

```
[4]: numberOfAirports = airportsDF.groupBy("country").agg({"id": "count"})
    airportsDF.createOrReplaceTempView("airportsDF")
```

```
[6]: numberOfAirportsSorted = spark.sql('''SELECT country, count(iata) as
    ↳ number_of_airports FROM airportsDF
                                     GROUP BY country Having count(iata) >
    ↳ 200 ORDER BY number_of_airports DESC''')
```

```
[7]: numberOfAirportsSorted.show()
```

```
[Stage 7:=====> (179 + 3) / 200]
```

```
+-----+-----+
```

| country       | number_of_airports |
|---------------|--------------------|
| United States | 1512               |
| Canada        | 430                |
| Australia     | 334                |
| Russia        | 264                |
| Brazil        | 264                |
| Germany       | 249                |
| China         | 241                |
| France        | 217                |

```
[ ]: # Find the top-10 airlines by total number of flights
```

```
[5]: numberOfFlights = flightsDF.groupBy("airline_id").agg({"airline_id": "count"})
numberOfFlights = numberOfFlights.withColumnRenamed("count(airline_id)", "num_flights")
joinedAirlineFlights = airlinesDF.join(numberOfFlights, airlinesDF.airline_id == numberOfFlights.airline_id)
icaoAirline = joinedAirlineFlights.selectExpr("name", "icao", "num_flights")
icaoAirlineOrdered = icaoAirline.sort("num_flights", ascending=False)
icaoAirlineOrderedFinal = icaoAirlineOrdered.limit(10).sort("num_flights", ascending=True)
```

```
[153]: icaoAirlineOrderedFinal.show(truncate=False)
```

```
[Stage 351:=====>(199 + 1) / 200]
```

| name                    | icao | num_flights |
|-------------------------|------|-------------|
| easyJet                 | EZY  | 1130        |
| Southwest Airlines      | SWA  | 1146        |
| Air China               | CCA  | 1260        |
| China Eastern Airlines  | CES  | 1263        |
| China Southern Airlines | CSN  | 1454        |
| US Airways              | USA  | 1960        |
| Delta Air Lines         | DAL  | 1981        |
| United Airlines         | UAL  | 2180        |
| American Airlines       | AAL  | 2354        |
| Ryanair                 | RYR  | 2484        |

```
[ ]: # Find the top-5 routes in the world
```

```
[6]: groupedFlights = flightsDF.groupBy("airport_source_id",  
    ↪ "airport_destination_id").agg({"airport_source_id": "count"})
```

```
[7]: groupedFlightsOrdere = groupedFlights.sort("count(airport_source_id)",  
    ↪ ascending=False)
```

```
[156]: groupedFlightsOrdere.count()
```

```
[156]: 37505
```

```
[8]: from pyspark.sql.functions import col  
joined1 = groupedFlightsOrdere.join(airportsDF, groupedFlightsOrdere.  
    ↪ airport_source_id == airportsDF.id).select(col("airport_source_id"),  
    col("airport_destination_id"), col("count(airport_source_id)").  
    ↪ alias("route_count"), col("name").alias("source_airport"),  
    col("city").alias("source_city"))
```

```
[9]: joined2 = joined1.join(airportsDF, joined1.airport_destination_id == airportsDF.  
    ↪ id).select(  
    col("source_airport"), col("source_city"), col("name").  
    ↪ alias("destination_airport"), col("city").alias("destination_city"),  
    ↪ col("route_count"))
```

```
[10]: routes = joined1.join(airportsDF, joined1.airport_destination_id == airportsDF.  
    ↪ id).select(col("airport_source_id"), col("airport_destination_id"))
```

```
[11]: routesName = joined2.sort("route_count", ascending=False)
```

```
[41]: routesName.show(5, truncate=False)
```

```
[Stage 53:=====> (8 + 1) / 9]
```

```
+-----+-----+-----+  
-----+-----+-----+  
|source_airport  
|source_city|destination_airport  
|destination_city|route_count|  
+-----+-----+-----+  
-----+-----+-----+  
|Chicago O'Hare International Airport          |Chicago |Hartsfield Jackson  
Atlanta International Airport |Atlanta |20 |  
|Hartsfield Jackson Atlanta International Airport|Atlanta |Chicago O'Hare  
International Airport          |Chicago |19 |  
|Phuket International Airport          |Phuket |Suvarnabhumi
```

|  |             |         |                 |
|--|-------------|---------|-----------------|
| Airport  | Bangkok     | 13      |                 |
| Chicago O'Hare International Airport             |             | Chicago | Louis Armstrong |
| New Orleans International Airport                | New Orleans | 13      |                 |
| Hartsfield Jackson Atlanta International Airport | Atlanta     | Miami   |                 |
| International Airport                            | Miami       | 12      |                 |

only showing top 5 rows

```
[ ]: # Create the graph of flight connections
# Build a graph using GraphFrames where vertices are the airports in airports.
↪ csv,
# and edges are the flights from one airport to another contained in routes.csv
```

```
[ ]: # Cleaning the flights (routes) by removing missing values
```

```
[142]: flightsDF.count()
```

```
[142]: 67663
```

```
[10]: flightsClean = flightsDF.filter((flightsDF["airport_source_id"] != "\\N") &
↪ (flightsDF["airport_destination_id"] != "\\N"))
```

```
[18]: flightsClean.count()
```

```
[18]: 67240
```

```
[ ]: from graphframes import GraphFrame
```

```
[ ]: v = airportsDF.select("id", "name", "icao", "city", "country")
v = v.withColumn("id", v.id.cast("string"))
```

```
[ ]: e = flightsClean.selectExpr("airport_source_id as src", "airport_destination_id
↪ as dst", "airline_id")
e = e.withColumn("src", e.src.cast("string"))
e = e.withColumn("dst", e.dst.cast("string"))
```

```
[ ]: g = GraphFrame(v, e)
```

```
[ ]: # Find the 5 airports with the highest ratio of outgoing edges over incoming
      ↪ edges (edgeOut/edgeIn)
```

```
[ ]: outgoing = g.outDegrees
      incoming = g.inDegrees
```

```
[ ]: outgoing_incoming = outgoing.join(incoming, outgoing.id == incoming.id).select(
      outgoing["id"], incoming["inDegree"], outgoing["outDegree"],
      ↪(outgoing["outDegree"] / incoming["inDegree"]).alias("ratio"))
```

```
[ ]: outgoing_incoming_final = outgoing_incoming.join(airportsDF, outgoing_incoming.
      ↪id == airportsDF.id).select(
      airportsDF["name"], outgoing_incoming["id"], outgoing_incoming["inDegree"],
      ↪outgoing_incoming["outDegree"], outgoing_incoming["ratio"])
```

```
[156]: outgoing_incoming_final.sort("ratio", ascending=False).show(5, truncate=False)
```

[Stage 493:=====> (163 + 1) / 200]

| name   | id   | inDegree | outDegree | ratio |
|--|------|----------|-----------|-------|
| Bunia Airport                                  | 1033 | 1        | 3         | 3.0   |
| Pikangikum Airport                             | 5521 | 2        | 5         | 2.5   |
| Transilvania Târgu Mureş International Airport | 1662 | 2        | 5         | 2.5   |
| Bahir Dar Airport                              | 1111 | 1        | 2         | 2.0   |
| Ivalo Airport                                  | 428  | 1        | 2         | 2.0   |

only showing top 5 rows

---

```
[ ]: # Finding the number of airports that from there we can reach the city of
      ↪ "Torino" taking exactly 1 flight
```

```
[ ]: # Finding the id of 'Turin Airport'
```

```
[27]: idTurin = outgoing_incoming_final.filter("name = 'Turin Airport'")
```

```
[8]: idTurin.show(1)
```

```
[ ]: # Create the motif
```

```
[159]: motifs = g.find("(a)-[]->(b)")
```

```
[ ]: # filter all destination with 'Turin Airport' id
```

```
[160]: airportsToTurin = motifs.filter("b.id = 1526")
```

```
[162]: airportsToTurin.select("a").distinct().count()
```

```
[162]: 29
```

---

```
[ ]: # Find the number of airports that can be reached from the city of "Torino" ↵  
      ↪ taking exactly 1 flight
```

```
[163]: airportsFromTurin = motifs.filter("a.id = 1526")
```

```
[165]: airportsFromTurin.select("b").distinct().count()
```

```
[165]: 29
```

```
[ ]: '''  
      In the analysis of Turin Airport's network connectivity, both inDegrees and ↵  
      ↪ outDegrees are observed to be 29.  
      This symmetry implies a balanced exchange of flights - for each departing ↵  
      ↪ flight, there is a corresponding  
      incoming flight. The equal values of inDegrees and outDegrees indicate a ↵  
      ↪ reciprocity in the airport's connectivity,  
      highlighting a comprehensive and evenly distributed network.  
      '''
```

---

```
[ ]: # Find the number of airports that can be reached from the city of "Torino" ↵  
      ↪ taking exactly 2 flights
```

```
[ ]: # Create the motif
```

```
[166]: stopFlightMotifs = g.find("(v1)-[]->(v2);(v2)-[]->(v3)")
```

```
[ ]: # Filter the origin airport 'Turin Airport' id
```

```
[167]: stopFlightFromTurin = stopFlightMotifs.filter("v1.id = 1526")
```

```
[168]: stopFlightFromTurin.select("v3").distinct().count()
```

[168]: 590

---

```
[ ]: # Find the number of airports that from there we can reach "Los Angeles_
    ↪International Airport" using less hop
    # than to reach the city of "Torino"
```

```
[ ]: # Find the id of "Los Angeles International Airport"
```

```
[21]: idLA = outgoing_incoming_final.filter("name = 'Los Angeles International_
    ↪Airport'")
```

```
[9]: idLA.show(1)
```

```
[ ]: # Calculate shortestPaths to "Los Angeles International Airport"
```

```
[24]: toLA = g.shortestPaths(landmarks=["3484"])
```

```
[ ]: # Calculate shortestPaths to "Turin Airport"
```

```
[36]: toTurin = g.shortestPaths(landmarks=["1526"])
```

```
[ ]: # Join two previous dataframes
```

```
[38]: LA_TO = toLA.join(toTurin, toLA.id == toTurin.id).select(toLA["id"],
    ↪toLA["name"].alias("toLA_name"),
    toLA["distances"].alias("toLA_distance"), toTurin["name"].
    ↪alias("toTurin_name"), toTurin["distances"].alias("toTurin_distances"))
```

```
[63]: LA_TO = LA_TO.withColumn("toLA_distance_numeric", col("toLA_distance").
    ↪getItem("3484")).withColumn(
    "toTurin_distance_numeric", col("toTurin_distances").getItem("1526"))
```

```
[ ]: # Filter those airports with less hops than Torino
```

```
[64]: LA_TO_filtered = LA_TO.filter("toLA_distance_numeric <
    ↪toTurin_distance_numeric")
```

```
[65]: LA_TO_filtered.count()
```

[65]: 1831



```
[ ]: # Find the number of airports that from there we can reach the city of Torino  
    ↪ using less hops than to reach  
    # "Los Angeles International Airport"
```

```
[66]: TO_LA_filtered = LA_TO.filter("toLA_distance_numeric >_  
    ↪ toTurin_distance_numeric")
```

```
[67]: TO_LA_filtered.count()
```

[67]: 94

---

```
[ ]: # Find the number of airports that from there we can reach with the same number  
    ↪ of hops Torino  
    # and "Los Angeles International Airport"
```

```
[68]: TOLA_equal_filtered = LA_TO.filter("toLA_distance_numeric ==_  
    ↪ toTurin_distance_numeric")
```

```
[69]: TOLA_equal_filtered.count()
```

[69]: 1244

---

```
[ ]: # Find the number of connected components of at least two airports are there in  
    ↪ the graph  
    # and the size of those connected components.
```

```
[28]: sc.setCheckpointDir("tmp_ckpts")
```

23/12/26 09:57:03 WARN spark.SparkContext: Spark is not running in local mode,  
therefore the checkpoint directory must not be on the local filesystem.  
Directory 'tmp\_ckpts' appears to be on the local filesystem.

```
[30]: gClean = g.dropIsolatedVertices()  
connComp = gClean.connectedComponents()  
connComp.createOrReplaceTempView("connComp")  
connCompClean = spark.sql(''SELECT component, count(*) FROM connComp GROUP BY  
    ↪ component Having count(*) >= 2'')
```

```
[34]: connCompClean.distinct().count()
```

[34]: 7

```
[41]: connCompClean.show()
```

```
+-----+-----+
| component | count(1) |
+-----+-----+
0	3188
9	4
721554505735	2
266287972363	4
300647710722	10
352187318278	4
1460288880648	2
+-----+-----+
```

---

```
[ ]: # By considering only the subgraph of the flights that are performed by two
      ↪different airlines (identified by the
      # icao), each involving at least 5 cities, I am going to select two airlines to
      ↪show their graphs.
```

```
[ ]: # Select two Airlines with previous characteristics
      # XLA (Excel Airways), GLG (Aerolineas Galapagos)
```

```
[ ]: test_city = icao_joined.join(airportsDF, icao_joined.src == airportsDF.id).
      ↪select(icao_joined['icao'], airportsDF['city'])
```

```
[ ]: test_city.createOrReplaceTempView("test_city")
```

```
[ ]: select_city = spark.sql('''
                                SELECT icao, COUNT(city)
                                FROM test_city
                                GROUP BY icao
                                HAVING COUNT(city) >= 5''')
```

```
[ ]: # Create a new dataframe (graphframe) with desired columns
```

```
[16]: icao_joined = flightsClean.join(airlinesDF, flightsClean.airline_id ==
      ↪airlinesDF.airline_id).select(
      ↪    flightsClean["airport_source_id"].alias("src"),
      ↪    flightsClean["airport_destination_id"].alias("dst"),
      ↪    airlinesDF["airline_id"], airlinesDF["icao"].alias("icao"),
      ↪    airlinesDF["name"].alias("airline_name"))
```

```
[ ]: # Create the subgraph of selected airlines
```

```
[17]: icao_joined.createOrReplaceTempView("icao_joined")
e_icao = spark.sql('''SELECT src, dst, icao, airline_name FROM icao_joined
↳WHERE icao = "XLA" OR icao = "GLG"''')
v_cities = airportsDF.select("id", "city")
g_icao = GraphFrame(v_cities, e_icao)
g_final = g_icao.dropIsolatedVertices()
```

```
[ ]: # Plot the subgraph of these flights
```

```
[64]: from graphviz import Digraph
def vizGraph(edge_list,node_list):
    Gplot=Digraph()
    edges=edge_list.collect()
    nodes=node_list.collect()
    for row in edges:
        Gplot.edge(str(row['src']),str(row['dst']),label=str(row['icao']))
    for row in nodes:
        Gplot.node(str(row['id']),label=str(row['city']))
    return Gplot
```

```
[1]: Gplot=vizGraph(g_final.edges, g_final.vertices)
Gplot
# the picture of final graph will be added at the end of the report.
```

---

```
[ ]: '''Find the destination airport at minimum distance from "Tancredo Neves
↳International Airport"
that we can reach by taking exactly 2 flights.I am going to return the
↳destination airport and its distance
from "Tancredo Neves International Airport" by considering that we cannot come
↳back to the
Tancredo Neves International Airport'''
```

```
[ ]: # Find the id of "Tancredo Neves International Airport"
```

```
[76]: idTN = outgoing_incoming_final.filter("name = 'Tancredo Neves International
↳Airport'")
```

```
[2]: idTN.show(1)
```

```
[ ]: # Create the motif that we cannot comeback
```

```
[69]: motifTN = g.find("(a)-[]->(b);(b)-[]->(c); !(c)-[]->(a);!(b)-[]->(a)")
```

```

[70]: fromTN = motifTN.filter("a.id = 2537")

[163]: fromTN.count()

]]

[163]: 39

[71]: destinationList = fromTN.select("c.id")

[25]: lat_lon_TN = airportsDF.filter("name = 'Tancredo Neves International Airport'").
      ↪show(5)

[78]: lat_lon = destinationList.join(airportsDF, destinationList.id == airportsDF.id).
      ↪select(
          destinationList['id'].alias('id'), airportsDF['name'],
      ↪airportsDF['latitude'], airportsDF['longitude'])

[ ]: # Define a function (haversine) to calculate the distance between airports

[116]: import math
def haversine(lat1, lon1, lat, lon):
    R = 6371.0
    lat1, lon1, lat, lon = map(math.radians, [lat1, lon1, lat, lon])
    dlat = lat - lat1
    dlon = lon - lon1
    hav = math.sin(dlat / 2) ** 2 + math.cos(lat1) * math.cos(lat) * math.
    ↪sin(dlon / 2) ** 2
    distance = 2 * R * math.asin(math.sqrt(hav))
    return distance

[ ]: # Register the function

[105]: spark.udf.register('hav', haversine)

24/01/09 17:57:13 WARN analysis.SimpleFunctionRegistry: The function hav
replaced a previously registered function.

[105]: <function __main__.haversine(lat1, lon1, lat, lon)>

[106]: dinstanceAirportsDF = lat_lon.selectExpr("id", "name", "hav(-19.62444305419922,
      ↪-43.97194290161133, latitude, longitude) as distance").sort("distance")

[118]: dinstanceAirportsDF.show(1, truncate = False)

[Stage 596:=====>(198 + 2) / 200]

```

| id   | name                               | distance           |
|------|------------------------------------|--------------------|
| 2555 | Hercílio Luz International Airport | 1008.7177113996354 |

only showing top 1 row

```
[ ]: '''There are 39 airports that can be reach by taking exactly two flights from
      "Tancredo Neves International Airport" that there is no return flight from
      ↪those destination.
      The minimum distance from this airport with these conditions is 1008.
      ↪7177113996354 which the destination
      airport is "Hercílio Luz International Airport".'''
```

```
[ ]: # Compute the total flown distance in kilometers, considering the distance from
      # "Tancredo Neves International Airport" to the first airport summed to the
      ↪distance from the first airport to
      # second one (total flown distance)
```

```
[94]: bothAirportsList = fromTN.select("b.id", "c.id").filter("c.id = 2555")
```

```
[26]: bothAirportsList.show()
```

```
[98]: airportB_List = fromTN.select("b.id")
```

```
[101]: lat_lon_B = airportB_List.join(airportsDF, airportB_List.id == airportsDF.id).
      ↪select(
          airportB_List['id'].alias('id'), airportsDF['name'],
          ↪airportsDF['latitude'], airportsDF['longitude'])
```

```
[27]: lat_lon_B.filter("id = 2442").show(1, truncate=False)
```

```
[121]: airport_B_C = lat_lon.selectExpr("id", "name", "hav(-34.5592, -58.4156,
      ↪latitude, longitude) as distance").filter("id = 2555")
```

```
[122]: airport_B_C.show(truncate=False)
```

[Stage 686:=====>

(58 + 17) / 75]

| id   | name                 | distance           |
|------|----------------------|--------------------|
| 2555 | Hercílio Luz Inte... | 1210.6138281258875 |

```
+-----+-----+-----+
```

```
[119]: airport_A_B = lat_lon_B.selectExpr("id", "name", "hav(-19.62444305419922, -43.
↪97194290161133, latitude, longitude) as distance")
```

```
[120]: airport_A_B.show(1, truncate=False)
```

```
[Stage 632:=====> (7 + 1) / 8]
```

```
+-----+-----+-----+
|id  |name                |distance      |
+-----+-----+-----+
|2442|Jorge Newbery Airpark|2186.151457533313|
+-----+-----+-----+
```

only showing top 1 row

```
[ ]: '''The distance between "Tancredo Neves International Airport" and "Jorge
↪Newbery Airpark" is 2186.151457533313
and the distance between "Jorge Newbery Airpark" and "Hercílio Luz
↪International Airport" is 1210.6138281258875,
so the sum of distances is 3396.76529.'''
```